

TMA4300 - Exercise 1

Stochastic simulation

Marcus A. Engbretnsen and Gina Magnussen

Problem A

Probability integral transform, rejection sampling and bivariate techniques

1. Sampling from $g(x)$ - Probability integral transform

The probability density function

$$g(x) = \begin{cases} cx^{\alpha-1}, & 0 < x < 1 \\ ce^{-x} & 1 \leq x \\ 0 & \text{otherwise,} \end{cases}$$

is given with c as a normalising constant and $\alpha \in (0, 1)$.

a)

The cumulative distribution function $G_X(x)$ is then found by integrating over the different domains and taking into account that the cdf must be continuous,

$$G_X(x) = \begin{cases} \frac{c}{\alpha} & 0 < x < 1, \\ \frac{c}{\alpha} + ce^{-1} - ce^{-x} & 1 \leq x \\ 0 & \text{else} \end{cases}$$

By using the property that the area under a pdf should integrate to 1, we find that $c = \frac{e\alpha}{e+\alpha}$

The probability integral transform, setting $u = G_X(x)$ and solving for x , gives

$$\begin{aligned} x = G^{-1}(u) &= \left(u \frac{e+\alpha}{e}\right)^{\frac{1}{\alpha}}, & 0 < x < 1 \\ x = G^{-1}(u) &= -\log \left[\left(\frac{e+\alpha}{e\alpha}\right) - \frac{u}{c} \right] = -\log \left[\frac{1}{c}(1-u) \right] & 1 \leq x \end{aligned}$$

b)

R function

R plot: gsample

2. Gamma distribution with $\alpha \in (0, 1)$, $\beta = 1$ generated by rejection sampling

$$f(x) = \begin{cases} \frac{1}{\Gamma(\alpha)} x^{\alpha-1} e^{-x}, & 0 < x \\ 0 & \text{else} \end{cases}$$

with $\alpha \in (0, 1)$ and $\beta = 1$.

a)

The acceptance probability R :

$$\frac{1}{\gamma} \frac{f(x)}{g(x)}$$

$$\gamma_1 = \frac{f(x)}{g(x)} = \begin{cases} \frac{1}{\Gamma(\alpha)} \frac{e^{-x}}{c} & 0 < x < 1 \\ \frac{1}{\Gamma(\alpha)} \frac{x^{\alpha-1}}{c} & x \geq 1 \end{cases}$$

$$\frac{\partial \gamma_1}{\partial x} = \begin{cases} -\frac{1}{\Gamma(\alpha)} \frac{e^{-x}}{c} & 0 < x < 1 \\ \frac{\alpha-1}{\Gamma(\alpha)} \frac{x^{\alpha-2}}{c} & x \geq 1 \end{cases}$$

Solving $\frac{\partial \gamma_1}{\partial x} = 0$ for x and inserting into the expression for R will now maximize the acceptance probability given the constraint that $f(x) < \frac{1}{\gamma} g(x)$. CHECK CONSTANTS!!

$$\gamma_1 = \begin{cases} \frac{1}{\Gamma(\alpha)} \frac{1}{c} & 0 < x < 1 \\ \frac{1}{\Gamma(\alpha)} \frac{1}{c} & x \geq 1 \end{cases}$$

which in turn gives an acceptance probability

$$R = \begin{cases} e^{-x} & 0 < x < 1 \\ x^{\alpha-1} & x \geq 1 \end{cases}$$

b) R function

3. Ratio of uniforms - gamma, $\alpha > 1, \beta = 1$.

Consider now the same distribution, but with parameters $\alpha > 1$ and $\beta = 1$. The ratio of uniforms method is used to simulate samples from this distribution. With C_f , f^* , a , b_+ and b_- as given in formula (3) and (4) in the exercise description, we find that

Need to find maximum of $f^*(x)$ to find the bounds for the area C_f . Since $f^*(x)$ is a concave function, we solve

$$\frac{df^*(x)}{dx} = (\alpha - 1)x^{\alpha-2} - x^{\alpha-1}e^{-x} = 0 \implies x = \alpha - 1 \frac{d(x^2 f^*(x))}{dx} = 0 \implies x = \alpha + 1$$

$$f^*(\alpha - 1) = (\alpha - 1)^{\alpha-1} e^{1-\alpha}$$

$$a = \sqrt{\sup_x f^*(x)} = \sqrt{f^*(\alpha - 1)} = \sqrt{(\alpha - 1)^{\alpha-1} e^{1-\alpha}}$$

$$b_+ = \sqrt{\sup_{x \geq 0} x^2 f^*(x)} = \sqrt{(\alpha + 1)^2 f^*(\alpha + 1)} = (\alpha + 1)^{\frac{\alpha+1}{2}} e^{-\frac{\alpha+1}{2}}$$

$$b_- = (\sqrt{\sup_{x \leq 0} x^2 f^*(x)}) = 0$$

$$C_f = [0, a] \times [b_-, b_+] = \left[0, \sqrt{(\alpha - 1)^{\alpha-1} e^{1-\alpha}}\right] \times \left[0, (\alpha + 1)^{\frac{\alpha+1}{2}} e^{-\frac{\alpha+1}{2}}\right]$$

Ratio of uniforms:

$$x_1 = \frac{u_2}{u_1} \qquad x_2 = u_1$$

Problem B

The Dirichlet distribution: Simulating using known relations

$$\begin{aligned} f_z(z, \alpha) &= dz_1 \dots dz_k \propto (z_1^{\alpha_1-1}) e^{-z_1} \dots (z_k^{\alpha_k-1}) e^{-z_k} dz_1 \dots dz_k \\ &= z_1^{\alpha_1-1} \dots z_k^{\alpha_k-1} e^{-(z_1 + \dots + z_k)} dz_1 \dots dz_k \\ &= z_1^{\alpha_1-1} \dots z_k^{\alpha_k-1} e^{-v} dz_1 \dots dz_k \end{aligned}$$

where $v = -(z_1 + \dots + z_k)$

Change of variables

$$z_i = x_i \cdot v \implies dz_i = dx_i \cdot v + x_i dv$$

Using $\sum_{i=1}^k x_i = 1$ and $dx_1 + \dots + dx_k = 0$ Define $w = dz_1 + \dots + dz_{k-1} = [dx_1 + \dots + dx_{k-1}]v + [x_1 + \dots + x_k]dv$

Then

$$\begin{aligned} dz_k &= dx_k v + x_k dv \\ &= -[dx_1 + \dots + dx_{k-1}]v + [1 - [x_1 + \dots + x_{k-1}]]dv \\ &= dv - ([dx_1 + \dots + dx_{k-1}]v + [x_1 + \dots + x_{k-1}]dv) \\ &= dv - w \end{aligned}$$

Using exterior algebra:

$$\begin{aligned} dz_1 \wedge \dots \wedge dz_{k-1} \wedge dz_k &= (dz_1 \wedge \dots \wedge dz_{k-1}) \wedge (dv - w) \\ &= \dots \\ &= v^{k-1} dx_1 \wedge \dots \wedge dx_{k-1} \wedge dv \end{aligned}$$

Filling into the expression gives:

$$\begin{aligned} f_z(z, \alpha) dz_1 \dots dz_k &\propto (x_1 v)^{\alpha_1-1} \dots (x_{k-1} v)^{\alpha_{k-1}-1} (v(1 - [x_1 + \dots + x_{k-1}]))^{\alpha_k-1} e^{-v^{k-1}} dx_1 \wedge \dots \wedge dx_{k-1} \wedge dv \\ &= v^{\alpha_1 + \dots + \alpha_{k-1}} e^{-v} dv \left(x_1^{\alpha_1-1} \dots x_{k-1}^{\alpha_{k-1}-1} \right) \left(1 - \sum_{i=1}^{k-1} x_i \right)^{\alpha_k-1} dx_1 \dots dx_{k-1} dv \end{aligned}$$

Problem C

A toy Bayesian model: Birthdays

The probability that two or more students has their birthday on the same day can be simulated as follows: Randomly assign a birth date to the 35 in a given NTNU class, and check for duplicate dates. If so, count this event as a success. If not, assign dates randomly again and to the same check. Do this many times. The estimate is then the number of successes divided by the number of trials.

1: Independent birthdays, equally likely

a)

```

# Birthdays
sim <- 10000
stud <- 35
count <- 0

for (i in 1:sim){
  bdays <- round(runif(stud)*365)
  if (sum(duplicated(bdays))>=1){
    count <- count + 1
  }
}

prob <- count/sim
print(prob)

```

b)

Exact calculation: The probability that two or more students has their birthday on the same day, is the complement of the event that no students has their birthday on the same day. Thus

$$\begin{aligned}
 P(\text{no. of students with same birthday} \geq 2) &= 1 - P(\text{no students with same birthday}) \\
 &= 1 - \frac{365!}{330! 365^{35}} = 0.8144
 \end{aligned}$$

The difference between the simulated and exact answer is of order 10^{-3} or better with enough simulations, meaning the simulated probability is good.

2: Bayesian model

a)

The posterior distribution of (q_1, q_2, q_3, q_4) , the distribution of the probabilities of being born in a specific season given observations, is


$$\begin{aligned}
 P(q_{1:4}|x_{1:4}) &= \frac{P(x_{1:4}|q_{1:4})P(q_{1:4})}{\text{integral of something } P(x_{1:4})} \\
 &\propto P(x_{1:4}|q_{1:4})P(q_{1:4}) \\
 &\propto \prod_{i=1}^k q_i^{x_i} \cdot \prod_{i=1}^k q_i^{\alpha_i-1} \\
 &\propto \prod_{i=1}^k q_i^{x_i+\alpha_i-1} \\
 &\propto \prod_{i=1}^k q_i^{\bar{\alpha}_i-1}
 \end{aligned}$$

where $(x_{1:4}) = (x_1, x_2, x_3, x_4)$ and similary for q . Thus, the posterior distribution of (q_1, q_2, q_3, q_4) is a Dirichlet distribution with parameters $\bar{\alpha}_i = x_i + \alpha_i$, $i = 1, \dots, 4$. - What is k?

b)

When the joint posterior distribution of $q_{1:4}$ is Dirichlet, the marginal distribution for each season is $q_i \sim \text{beta}\left(\alpha_i, \sum_{k=1}^4 \alpha_k - \alpha_i\right)$.

These marginal distributions are plotted by simulating several Dirichlet distributions $x = (x_1, \dots, x_K)$, extracting for instance all x_1 s, which are beta distributed with the same parameters, and making a histogram of these samples. The built-in function in R is used to verify that the simulation of the beta distributions, and thus also the Dirichlet distribution, is correct. ?? Built-in or sample? Another option is to simulate the marginal beta distributions by using the relation between the gamma and the beta distribution, since we in the previous tasks have shown this sample generator to be reasonably correct.



Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

3: Sample/Estimate p

Given the probabilities q_i , of being born within a specific season we want to find the probability of p . This is here done by simulation. Denote

$$N_j = \text{no. of students born in season } j, \quad j \in \{1 = \text{spring}, \dots, 4 = \text{winter}\}$$

and set $m = 35$ again.

a) Sampling from posterior distribution $P(q_1, q_2, q_3, q_4 | x_1, x_2, x_3, x_4)$ and $P(N_1, N_2, N_3, N_4 | q_1, q_2, q_3, q_4)$

Sampling from posterior distribution $P(q_1, q_2, q_3, q_4 | x_1, x_2, x_3, x_4)$

KODE

As in C.2, the distribution of $(N_1, N_2, N_3, N_4 | q_1, q_2, q_3, q_4)$ will be multinomial with parameters (m, q_1, q_2, q_3, q_4) , but with $m = 35$.

$$(q_1, q_2, q_3, q_4) \sim \text{Dirichlet}(\alpha_1, \alpha_2, \alpha_3, \alpha_4)$$

$$(N_1, N_2, N_3, N_4) \sim \text{multinomial}(m, q_1, q_2, q_3, q_4)$$

How to sample from a multinomial distribution: Divide the unit interval $[0, 1]$ into four subsequent intervals $[0, q_1] \cup (q_1, q_1 + q_2] \cup (q_1 + q_2, q_1 + q_2 + q_3] \cup (q_1 + q_2 + q_3, q_1 + q_2 + q_3 + q_4]$ with length q_1, q_2, q_3, q_4 respectively. Then sample $u \sim \text{Unif}[0, 1]$ and increment N_j in the corresponding interval, in this case season, for $j = 1, \dots, 4$. That is, if u is between q_1 and $q_1 + q_2$, then N_2 is incremented. Do this $m = 35$ times.

b)

A formula for

$$p = f(N_1, N_2, N_3, N_4)$$

can be found by seeing that

$$p = P[(Y_1 \cup Y_2 \cup Y_3 \cup Y_4) = 1] = 1 - P[(Y_1 \cup Y_2 \cup Y_3 \cup Y_4) = 0]$$

where Y_j is the event that two or more students have the same birthday in season j .

Since birthdays between different students are independent of each other, we can assume that the event of no students having the same birthday in the spring is independent of the event that no students have the same birthday in the winter, since the events are disjoint. So

$$p = 1 - [P(Y_1 = 0) + P(Y_2 = 0) + P(Y_3 = 0) + P(Y_4 = 0)]$$

Now, since the outcome of Y_i is binary, we can render the above equation to a function of N_i using that K_i is the number of days in each season. This gives

$$P(Y_i = 0) = \frac{K_i(K_i - 1)(K_i - 2) \dots (K_i - (N_i - 1))}{K_i!}$$

$$\binom{K_i}{N_i} \cdot \frac{N_i!}{K_i^{N_i}}$$

and so

$$p = 1 - \left[\sum_{i=1}^4 \binom{K_i}{N_i} \cdot \frac{N_i!}{K_i^{N_i}} \right]$$

Now p can be calculated by sampling from $(N_1, N_2, N_3, N_4 | q_1, q_2, q_3, q_4)$, the multinomial distribution, calculating p by using the formula above, repeating this n times and then calculate the mean.

Confidence interval of posterior mean of p : Sort all the n found values for p , and then remove the 2.5% highest and lowest values.

Posterior mean is larger than the probability in C.1. This is reasonable because the probabilities q determine how likely it is to have a birthday within one season. If one q_i is larger than another, there are more chances of birthdays colliding within that season, and thus the probability of two or more people having the same birthday increases.

c)

Assume now a Dirichlet prior with $\alpha_1 = \alpha_2 = \alpha_3 = \alpha_4 = 20$.

Importance sampling

Comment and explain result: - Small alphas, all equal, will not affect the distribution that much, the observations have a lot to say - Bigger alphas, if all are equal, will give uniformly distributed birthdays between seasons, and push the q_i s closer to 0.25 (since the observations become less important.) Meaning we should get a similar result as in C1. α_i significantly bigger than the others alphas will increase the q_i (seen in pdf expression for $q|x$), the chance of colliding birthdays is bigger and thus p increases.